

Recently we have received many GPU out of virtual memory issues from some sepecific using case in Gallery, Camera, Wechat, Browser/Webview based app, or just in general monkey/aging test.

As a result, it may cause the incorrect rendering like black blocks, or native crash then app process exits abnormally.

### Classic errors in logs:

Logcat log:

#### <snippet 1>

```
08-06 14:42:11.158 7553 9292 W Adreno-GSL: <sharedmem_gpuobj_alloc:1907>: sharedmem_gpumem_alloc: mmap failed errno 12
Out of memory
08-06 14:42:11.195 7553 9292 W Adreno-GSL: <sharedmem_gpuobj_alloc:1907>: sharedmem_gpumem_alloc: mmap failed errno 12
Out of memory
08-06 14:42:11.289 7553 9292 W Adreno-GSL: <sharedmem_gpuobj_alloc:1907>: sharedmem_gpumem_alloc: mmap failed errno 12
Out of memory
```

#### <snippet 2>

```
11-18 12:27:39.485 6414 6444 W Adreno-GSL: <sharedmem_gpuobj_alloc:2021>: sharedmem_gpumem_alloc: mmap failed errno 12
Out of memory
11-18 12:27:39.496 6414 6444 E Adreno-GSL: <gsl_memory_alloc_pure:2125>: GSL MEM ERROR: kgsi_sharedmem_alloc ioctl failed.
11-18 12:27:39.526 6414 6444 W Adreno-GSL: <sharedmem_gpuobj_alloc:2021>: sharedmem_gpumem_alloc: mmap failed errno 12
Out of memory
11-18 12:27:39.536 6414 6444 E Adreno-GSL: <gsl_memory_alloc_pure:2125>: GSL MEM ERROR: kgsi_sharedmem_alloc ioctl failed.
11-18 12:27:39.569 6414 6444 W Adreno-GSL: <sharedmem_gpuobj_alloc:2021>: sharedmem_gpumem_alloc: mmap failed errno 12
Out of memory
11-18 12:27:39.579 6414 6444 E Adreno-GSL: <gsl_memory_alloc_pure:2125>: GSL MEM ERROR: kgsi_sharedmem_alloc ioctl failed.
```

Kernel log:

#### <snippet 1>

```
<3>[ 1333.877571] (1)[9292:Thread-889]kgsi kgsi-3d0: |kgsi_get_unmapped_area| _get_svm_area: pid 7553 addr 3ab87000 pgoff 188
len 32833536 failed error -12
<3>[ 1333.915015] (1)[9292:Thread-889]kgsi kgsi-3d0: |kgsi_get_unmapped_area| _get_svm_area: pid 7553 addr 3ab87000 pgoff 188
len 32833536 failed error -12
<3>[ 1334.008835] (1)[9292:Thread-889]kgsi kgsi-3d0: |kgsi_get_unmapped_area| _get_svm_area: pid 7553 addr 3a6ed000 pgoff 185
len 32833536 failed error -12
```

#### <snippet 2>

```
11-18 12:27:39.632 0 0 I Kernel : <3>[76537.821739] kgsi kgsi-3d0: |kgsi_get_unmapped_area| _get_svm_area: pid 6414 addr 0 pgoff
11 len 8458240 failed error -12
11-18 12:27:39.632 0 0 I Kernel : <3>[76537.821739] kgsi kgsi-3d0: |kgsi_get_unmapped_area| _get_svm_area: pid 6414 addr 0 pgoff
11 len 8458240 failed error -12
```

11-18 12:27:39.632 0 0 | Kernel : <3>[76537.821739] kgsi kgsi-3d0: |kgsi\_get\_unmapped\_area| \_get\_svm\_area: pid 6414 addr 0 pgoff 11 len 8458240 failed error -12

### <snippet 3>

<3>[347733.422879]@6[11-11 15:55:59]kgsi kgsi-3d0: |kgsi\_get\_unmapped\_area| \_get\_svm\_area: pid 18294 addr 0 pgoff 88 len 8454144 failed error -12 SUBSYSTEM=kgsi DEVICE=c239:0

<3>[347733.429543]@6[11-11 15:55:59]kgsi kgsi-3d0: |kgsi\_get\_unmapped\_area| \_get\_svm\_area: pid 18294 addr 0 pgoff 5c len 8454144 failed error -12 SUBSYSTEM=kgsi DEVICE=c239:0

<3>[347733.435406]@5[11-11 15:55:59]kgsi kgsi-3d0: |kgsi\_get\_unmapped\_area| \_get\_svm\_area: pid 18294 addr 0 pgoff 5c len 8454144 failed error -12 SUBSYSTEM=kgsi DEVICE=c239:0

### Background knowledge:

1. In recent Gfx/kgsi driver, it enables use\_cpu\_map feature, so it will use the same virtual mapping on CPU and GPU. When user-mode call kgsi to alloc Gfx/gpu memory, the returned GPU address will be 0, then user-mode calls mmap() into kgsi driver again to get the valid GPU address.
2. When user-mode calls mmap(), it will calls to kgsi driver kgsi\_get\_unmapped\_area()->\_get\_svm\_area() to lookup the whole process virtual address space to find one unmapped continuous range of VMA, which can meet both CPU's mapping restriction and GPU's mapping restriction.  
In CPU side, the mapping restrictions are the Dalvik VM and Java Heap will occupy a huge range of VMA, also the same with dynamically loadable librarys (.so files), also other restrictions like stack, heap, page guards.  
In GPU side, there are similiar restrictions, GPU kgsi global memstore must be mapped directly, use fixed map; ION memory does not implement use CPU map, so only mapped to GPU; Secure buffer cannot be mapped to CPU, only mapped to GPU.
3. In long time running, due to frequently alloc/free memory and map/unmap virtual address, process's virtual address space may be badly fragmented. In such situation, even the total free space are big enough, but they are fragmented into small pieces, not continous, but mmap() always need a continous range free VMA in virtual address space to map. so it is very easy to cause mmap() failed for some big size map request.
4. You can read more useful info from Appendix - kgsi mmap virtual address range limitaion.

### Root cause analysis:

The root cause is user-mode mmap() is failed due to kernel/kgsi driver cannot find one big enough continous range free VMA in process's virtual address space.

Here are 2 possible reasons:

- #1. There is obvious memory leak in APP, which consumes too much range of vma, then cause virtual address space is used out. This is definitely app own issue, need to fix in app side.
- #2. There is no obvious memory leak in APP, but the virtual address space is fragmented badly, so cannot find a big enough continuous range VMA for a new request in mmap().

For #2, recently we have a known optimization fix of CR#1049887 in GFX adreno libs, which can improve the virtual address fragmentation introduced by GFX adreno libs.

In kernel log, if you see the addr parameter is not 0 like below, it indicate the known optimization is missed, so you can apply

the CR#1049887 fix at first.

kgs1 kgs1-3d0: |kgs1\_get\_unmapped\_area| \_get\_svm\_area: pid 7553 **addr 3ab87000** pgoff 188 len 32833536 failed error -12

**Just remind** CR#1049887 fix is only optimization in GFX adreno lib, cannot solve all the fragmentation issue.

You may still meet the issue even with the CR#1049887 fix, like below log:

kgs1 kgs1-3d0: |kgs1\_get\_unmapped\_area| \_get\_svm\_area: pid 6414 **addr 0** pgoff 11 len 8458240 failed error -12

kgs1 kgs1-3d0: |kgs1\_get\_unmapped\_area| \_get\_svm\_area: pid 18294 **addr 0** pgoff 88 len 8454144 failed error -12

## How to debug?

To debug the issue for both #1 and #2 reasons, the problem process's proc\_maps and kgs1\_mem are very valued info to analyze, especially the last time log before process crash.

You can apply below 2 linux shell scripts to get the continuous proc\_maps info and kgs1\_mem info when reproducing issue.

### get\_kgs1\_mem\_info.sh >

```
adb root
adb wait-for-device
adb shell "echo 0>/proc/sys/kernel/kptr_restrict"
while true;
do
adb shell echo "*****/d/kgs1/proc/<pid>/mem start *****">>kgs1_mem_info.txt
adb shell date>>kgs1_mem_info.txt
adb shell "cat /d/kgs1/proc/<pid>/mem">>kgs1_mem_info.txt
adb shell echo "*****/d/kgs1/proc/<pid>/mem completed *****">>kgs1_mem_info.txt
sleep 1
done
```

### get\_proc\_maps\_info.sh >

```
adb root
adb wait-for-device
adb shell "echo 0>/proc/sys/kernel/kptr_restrict"
while true;
do
adb shell echo "*****/proc/<pid>/maps start *****">>maps_info.txt
adb shell date>>maps_info.txt
adb shell echo "startaddr-endaddr flags offset major:minor size node">>maps_info.txt
adb shell "cat /proc/<pid>/maps">>maps_info.txt
adb shell echo "*****/proc/<pid>/maps completed *****">>maps_info.txt
sleep 1
done
```

**[note]** you must replace "<pid>" with the correct value of the problem process pid

When the issue is reproduced, just provide the complete kgsi\_mem\_info.txt and maps\_info.txt to QC GFX team to check.

## How to analyze the log?

We have 2 python script to analyze the kgsi\_mem info and proc\_maps info directly, then get the statistics data in a simple view format.

**check\_kgsi\_map.py** - output the kgsi mem info statistics data, like below

	Num (count)	Max (KB)	Sum (KB)
-----			
any(0)	40	256	1420
arraybuffer	17	4096	9648
command	251	128	5976
egl_surface	6	8220	24684
gl	227	768	4956
renderbuffer	1	1904	1904
texture	1147	4160	119808
-----			
Total	168396		

If we can see some abnormal state of GPU mem usage in the kgsi mem info, then it is APP memory leak issue.

In the above example info, we can see there are too many GL Texture memory, totally 1147 textures, the maximum texture is 4M+, and totally texture memory is 119M~.

We may suspect there is GL Texture leak in APP side, then we can ask for other logs, like GL API log to confirm this.

**check\_proc\_maps.py** - output the process maps info statistics data, like below

	Num (count)	Max (KB)	Sum (KB)
-----			
/dev/kgsi-3d0	1687	4164	150488
[anon:libc_malloc]	228	21504	354816
<~snip~>			
[anon:v8_heap]	26	1024	11756
[stack:11052]	1	1036	1036
[stack:11109]	1	1036	1036
<~snip~>			
[stack:9835]	1	1036	1036
[stack:9990]	1	1036	1036
[stack]	1	8188	8188
anon_inode:dmapuf	8	8220	25076
-----			
Total	1261092		

Hole Num	(count)	Max (KB)	Min (KB)	Sum (KB)
5545	135124	0	1218236	

If we can see too many entries in the proc maps info in specific usage, it may indicate some abnormal state in current process.

In the above example info, we can see there are totally 1687 entries of /dev/kgsl-3d0, the maximum size of the entry is 4M+, the total memory is 150M+, this aligned with the kgsl\_mem info log.

We can see 228 entries of libc\_malloc, the maximum size of the entry is 21M+, the total memory size is 354M+.

We can also see 211 entries of [stack:xxxx], where xxxx is standing for different number, each has only one entry, with fixed size 1036, it may be related to binder communication with different process.

These too many entries really cause the process virtual address space badly fragmented.

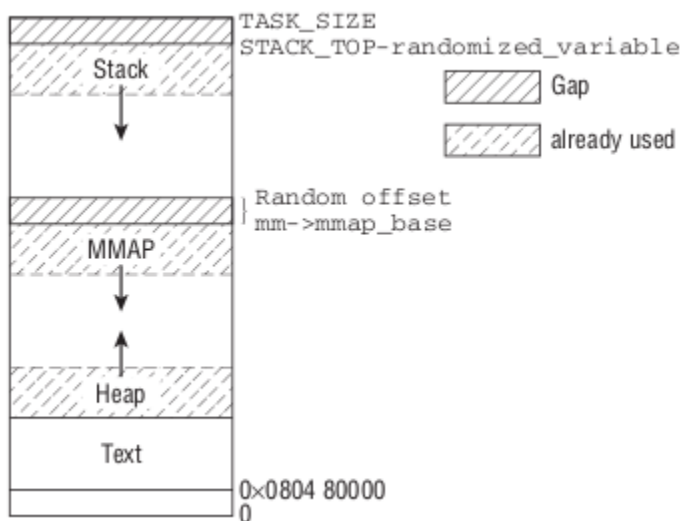
**[Note] These 2 python script (check\_kgsl\_map.py, check\_proc\_maps.py) are for Qualcomm internal debug only, cannot share to OEM customer.**

## Appendix - kgsl mmap virtual address range limitaion

1. As you know, the process user mode virtual address is managed by the mm\_struct \*mm in task\_struct, the virtual address layout is defined in mm\_struct \*mm as below, the user mode stack is started from STACK\_TOP, if PF\_RANDOMIZE is set, the start address will be reduced with a random value.

Each architecture must define STACK\_TOP, and most of them are defined as TASK\_SIZE, it is normally 0xC0000000 in 32bit device. After randomized, the process stack start address is save in mm->start\_stack, which can be read by cat /proc/xxx/stat.

As in the below figure, the stack is grown up from Top to Bottom. The range for mmap is starting from mm->mmap\_base, which is initialized by calling mmap\_base(), there is also one extra gap between stack and mmap range for isolation.



The mmap\_base() function is as below:

```

#define MIN_GAP (128*1024*1024)
#define MAX_GAP (TASK_SIZE/6*5)

static inline unsigned long mmap_base(struct mm_struct *mm)
{
    unsigned long gap = current->signal->rlim[RLIMIT_STACK].rlim_cur; // rlim_cur the default value is 8388608, which is 8M, can check it
    with getrlimit(RLIMIT_STACK, &limit)
    unsigned long random_factor = 0;
    if (current->flags & PF_RANDOMIZE)
        random_factor = get_random_int() % (1024*1024);

    if (gap < MIN_GAP) // MIN_GAP is used to guarantee the minimum stack size is at least 128M
        gap = MIN_GAP;
    else if (gap > MAX_GAP) // MAX_GAP the maximum stack size is TASK_SIZE/6*5, which is 2.5G
        gap = MAX_GAP;

    return PAGE_ALIGN(TASK_SIZE - gap - random_factor); // use the random_factor to avoid the the underflow of stack
}

```

The most recent code for android 8.0 is in /kernel/msm-4.9/arch/arm/mm/mmap.c

```

#ifdef CONFIG_HAVE_ARCH_MMAP_RND_COMPAT_BITS
const int mmap_rnd_compat_bits_min = CONFIG_ARCH_MMAP_RND_COMPAT_BITS_MIN;
const int mmap_rnd_compat_bits_max = CONFIG_ARCH_MMAP_RND_COMPAT_BITS_MAX;
int mmap_rnd_compat_bits __read_mostly = CONFIG_ARCH_MMAP_RND_COMPAT_BITS; // 16 <- 8, the original value is 8bits, the
range is 1M, the new value is 16bits, the range is 256M
#endif

/* gap between mmap and stack */
#define MIN_GAP (128*1024*1024UL)
#define MAX_GAP ((TASK_SIZE)/6*5)
static int mmap_is_legacy(void)
{
    if (current->personality & ADDR_COMPAT_LAYOUT)
        return 1;
    if (rlimit(RLIMIT_STACK) == RLIM_INFINITY)
        return 1;
    return sysctl_legacy_va_layout;
}

static unsigned long mmap_base(unsigned long rnd)
{
    unsigned long gap = rlimit(RLIMIT_STACK);
    if (gap < MIN_GAP)
        gap = MIN_GAP;
}

```

```

else if (gap > MAX_GAP)
gap = MAX_GAP;
return PAGE_ALIGN(TASK_SIZE - gap - rnd);
}
unsigned long arch_mmap_rnd(void)
{
unsigned long rnd;
rnd = get_random_long() & ((1UL << mmap_rnd_bits) - 1);
return rnd << PAGE_SHIFT;
}
void arch_pick_mmap_layout(struct mm_struct *mm)
{
unsigned long random_factor = 0UL;
if (current->flags & PF_RANDOMIZE)
random_factor = arch_mmap_rnd();
if (mmap_is_legacy()) {
mm->mmap_base = TASK_UNMAPPED_BASE + random_factor;
mm->get_unmapped_area = arch_get_unmapped_area;
} else {
mm->mmap_base = mmap_base(random_factor);
mm->get_unmapped_area = arch_get_unmapped_area_topdown;
}
}

```

2. In GPU kgsi driver, when user mode do a mmap(), it will call into kgsi driver function kgsi\_get\_unmapped\_area() and kgsi\_mmap().

In kgsi\_get\_unmapped\_area(), if it is not SECURE buffer, then we support use CPU map for SVM feature support, with this both CPU and GPU are using the exact same virtual address.

In \_get\_svm\_area() called from kgsi\_get\_unmapped\_area(), it will firstly get the svm range for GPU/kgsi side, then clamped the range with CPU's mmap requirements.

kernel/msm-4.9/drivers/gpu/msm/kgsi.c

```

static unsigned long _get_svm_area(struct kgsi_process_private *private,
struct kgsi_mem_entry *entry, unsigned long hint,
unsigned long len, unsigned long flags)
<~snippet~>
/* get the GPU pagetable's SVM range */
if (kgsi_mmu_svm_range(private->pagetable, &start, &end, // start is usually 0x00000000, end is usually 0xBF000000
entry->memdesc.flags))
return -ERANGE;
/* now clamp the range based on the CPU's requirements */
start = max_t(uint64_t, start, mmap_min_addr); // mmap_min_addr is usually 0x00008000

```

```
end = min_t(uint64_t, end, current->mm->mmap_base); // current->mm->mmap_base is calculated by the function  
if (start >= end)  
return -ERANGE;
```

so we must check the **start** and **end** paramters value here for the final valid virtual address range, which is highly affected by both GPU svm range and CPU mmap limit of current->mm->mmap\_base.

QUALCOMM®  
2018-01-29 18:46:58 PST  
miaoshoufei@meizu.com